

	Report	Result	Size	Time	Cycles	GPU	SM Frequency	Process	Attributes
 Current	28	528 - transposeKernelInt	(16, 16, 1)x(32, 32, 1)	53.86 us	54,810	0 - NVIDIA GeForce GTX 1650 with Max-Q Design	1.02 Ghz	[21182] main	
 Baseline 1	padding	528 - transposeKernelInt	(16, 16, 1)x(32, 32, 1)	25.41 us	25,681	0 - NVIDIA GeForce GTX 1650 with Max-Q Design	1.01 Ghz	[23857] main	

Summary

Details

Source

Context

Comments

Raw


Session


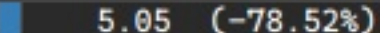
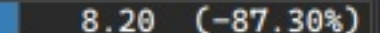

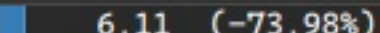
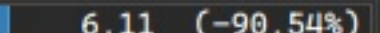
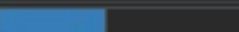

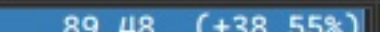

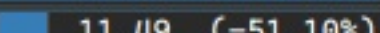
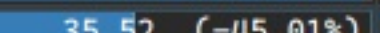

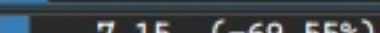
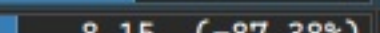

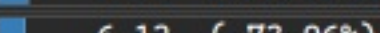
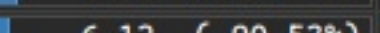

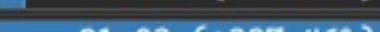
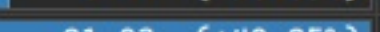
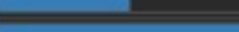



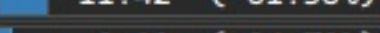
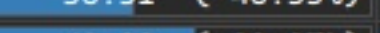

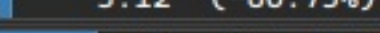

Compare

Tools

View

Export


This table shows all results in the report. Use the column headers to sort the results in this report. Double-click a result to see detailed metrics. Double-click on demangled names to rename it.

ID	Estimated Speedup	Function Name	Demangled Name	Duration (8.2008e+07)	Runtime Improvement (4.26926e+07)	Compute Throughput	Memory Throughput	# Registers	Grid Size
0	 52.47	checkSortedRowWi...	checkSortedRowWi...	0.05 (+81.99%)	0.02	 5.05 (-78.52%)	 8.20 (-87.30%)	16 (+0.00%)	16, 512,
1	 65.18	checkSortedColum...	checkSortedColum...	0.03 (+28.97%)	0.02	 6.11 (-73.98%)	 6.11 (-90.54%)	16 (+0.00%)	16, 512,
2	 23.28	sortRowsKernelInt	sortRowsKernelInt(...	22.63 (+88,966.75%)	5.27	 89.48 (+280.87%)	 89.48 (+38.55%)	18 (+12.50%)	512, 1,
3	 89.13	transposeKernelInt	transposeKernelInt(...	0.05 (+111.96%)	0.05	 11.49 (-51.10%)	 35.52 (-45.01%)	16 (+0.00%)	16, 16,
4	 52.18	checkSortedRowWi...	checkSortedRowWi...	0.08 (+198.99%)	0.04	 7.15 (-69.55%)	 8.15 (-87.38%)	16 (+0.00%)	16, 512,
5	 65.06	checkSortedColum...	checkSortedColum...	0.03 (+29.35%)	0.02	 6.12 (-73.96%)	 6.12 (-90.53%)	16 (+0.00%)	16, 512,
6	 28.37	sortRowsKernelInt	sortRowsKernelInt(...	22.25 (+87,467.76%)	6.31	 91.03 (+287.46%)	 91.03 (+40.95%)	18 (+12.50%)	512, 1,
7	 86.92	transposeKernelInt	transposeKernelInt(...	0.06 (+121.41%)	0.05	 11.42 (-51.38%)	 35.31 (-45.33%)	16 (+0.00%)	16, 16,
8	 87.19	checkSortedRowWi...	checkSortedRowWi...	32.41 (+127,439.17...	28.25	 3.12 (-86.73%)	 49.82 (-22.87%)	16 (+0.00%)	16, 512,
9	 59.99	checkSortedColum...	checkSortedColum...	4.43 (+17,318.14%)	2.66	 23.11 (-1.63%)	 40.01 (-38.05%)	16 (+0.00%)	16, 512,

Est. Speedup: 89.13%

locations. The [CUDA Best Practices Guide](#) has an example on optimizing shared memory accesses.

Shared Load Bank Conflicts

Est. Speedup: 68.81%

The memory access pattern for shared loads might not be optimal and causes on average a 32.0 - way bank conflict across all 8192 shared load requests.This results in 253952 bank conflicts, which represent 96.88% of the overall 262144 wavefronts for shared loads. Check the [Source Counters](#) section for uncoalesced shared loads.

Achieved Occupancy

Est. Speedup: 10.84%

The difference between calculated theoretical (100.0%) and measured achieved occupancy (89.2%) can be the result of warp scheduling overheads or workload imbalances during the kernel execution. Load imbalances can occur between warps within a block as well as across blocks of the same kernel. See the [CUDA Best Practices Guide](#) for more details on optimizing occupancy.